

# Sales Forecast with ARIMA

Forecast has been always a crucial and challenging task for a business. Given analysis is intended to explore and evaluate ARIMA forecasting techniques, its accuracy in prediction and application.

## 1. Profiling Data

Data set is derived prior to the analysis from a database and contains historical sales data for the department in interest.

Data set variables:

Yr - Year;

Mo\_no - Month;

TotalUniqShipTos - Monthly count of accounts invoiced;

TotalInv - Monthly sales invoices;

```
##      Yr Mo_No      Date TotalUniqShipTos  TotalInv
## 1  2021      2 2/1/2021             7215   4540000
## 2  2021      1 1/1/2021             7641 5059810.67
## 3  2020     12 12/1/2020             7006   4351108
## ...  ...    ...      <NA>             ...       ...
## 59 2016      4 4/1/2016             5403 2508213.51
## 60 2016      3 3/1/2016             6362 2926410.14
## 61 2016      2 2/1/2016             5753 2745165.31
## 62 2016      1 1/1/2016             5112 2442804.13
```

While exploring variable types, I see that Yr, Mo\_no and Date variables need to be transformed into a categorical data type:

```
##
## No. of observations = 62
## Variable      Class      Description
## 1 Yr           integer
## 2 Mo_No        integer
## 3 Date         character
## 4 TotalUniqShipTos integer
## 5 TotalInv     numeric
```

Next I sort data by date from latest to earliest. This is to assure that training set will consist of later months and validation set will contain the most recent months.

```
dataset <- dataset[order(dataset$Date, dataset$Mo_No), ]
```

### Missing Values Check:

Data set does not have any missing values, so I can move on.

```
any(is.na(dataset))
```

```
## [1] FALSE
```

Final look at the data.

```
attach(dataset)
des(dataset)
```

```
##
## No. of observations = 62
## Variable      Class      Description
## 1 Yr           factor
## 2 Mo_No        factor
## 3 Date          Date
## 4 TotalUniqShipTos integer
## 5 TotalInv      numeric
```

```
headTail(dataset, 6)
```

```
##      Yr Mo_No      Date TotalUniqShipTos  TotalInv
## 62  2016      1 2016-01-01           5112 2442804.13
## 61  2016      2 2016-02-01           5753 2745165.31
## 60  2016      3 2016-03-01           6362 2926410.14
## 59  2016      4 2016-04-01           5403 2508213.51
## 58  2016      5 2016-05-01           5295 2496541.34
## 57  2016      6 2016-06-01           6046 2810085.05
## ... <NA> <NA>      <NA>           ...      ...
## 4   2020     11 2020-11-01           7125  4459344
## 3   2020     12 2020-12-01           7006  4351108
## 2   2021      1 2021-01-01           7641 5059810.67
## 1   2021      2 2021-02-01           7215  4540000
```



## 2. Exploring Data Set

### Side Note:

The current data set consists of 5 years of data (2016-2020). Important to mention that April and May of 2020 were outliers caused by pandemic. Revenue during those months did not take its natural course of flow but rather were impacted by rare event. While synthesizing this data set, I replaced sales values for April and May with values derived from pre-pandemic forecast, which allowed me to still include the those months in my analysis and utilize them in the model.

Both Fig 1 and Fig 2 show the steady exponential trend of sales and served locations. However, from looking at the data I see that 2016 year is not quite in a line with the pattern observed in the later years. Considering required at least 3 years for training the model and latest 12 month (1 year) for validating the model, I decide to exclude 2016 year from this analysis and utilize records 13 to 61 that represent 2017-2020 years.

Going forward I *train* the model on previous 36 months (3 year) - Jan, 2017 to Jan, 2020 (inclusive) and *test* the model on last 12 months - Feb 2020 to Jan 2021 (inclusive).




### 3. Evaluating Time Series

Since this is a time series data, it is necessary to determine whether the data is **stationary**.


ACF plot shows the significant correlation between current time series and previous lags. Typically, such assumption is appropriate when time series, represented by vertical lines, cross the level of significance (blue dashed line), just as demonstrated below. Also, observation suggests presence of correlation between the last 3 terms, which assumes the lag equal to 3. Since time series drop gradually on ACF graph, there is an assumption of non-stationary data behavior. (Versus, if there was a sudden drop it would point on the stationary data and zero correlation between lags.)

PACF plot also suggests lag equal to 3, because of the first three vertical lines, that are above the 95% of significance level.

#### ACF for Sales



#### PACF for Sales



To verify above assumptions, `adf.test()` is performed. Null hypothesis is defined as data is non-stationary. Since p-value is more than assumed 0.05, we fail to reject null hypothesis and have enough evidence to assume that the data is non-stationary. `Adf.test()` also recommends lag of 3, which confirms above evaluation from the plots.

```
tseries::adf.test(TotalInv[13:61], alternative = "stationary")
```

```
##
## Augmented Dickey-Fuller Test
##
## data: TotalInv[13:61]
## Dickey-Fuller = -2.6275, Lag order = 3, p-value =
## 0.3231
## alternative hypothesis: stationary
```

#### 4. Fitting AUTO.ARIMA

Next I define the training data set as time series and apply `auto.arima()` algorithm to find fitted model that appear to be described as  $(0,1,1)(0,1,0)[12]$ .

Aiken value is 672.30, which sets the benchmark for choosing the best fitted model.

```
ts <- ts(dataset$TotalInv[13:49], frequency = 12)
ts1 <- window(ts)
(arima1 <- auto.arima(ts1, D = 1))

## Series: ts1
## ARIMA(0,1,1)(0,1,0)[12]
##
## Coefficients:
##          ma1
##       -0.8114
## s.e.    0.1198
##
## sigma^2 estimated as 7.243e+10: log likelihood=-334.15
## AIC=672.3   AICc=672.87   BIC=674.66
```


#### 5. Forecast of the next 12 months using AUTO.ARIMA(0,1,1)(0,1,0)[12] (Option 1)

By using `auto.arima()` in a **for loop** with iterator 12, I intend to forecast 12 month sales (validation set) by predicting one month at a time, while including previously predicted month in training set. By the time **for loop** is finished, I will have predicted sales for the 12 months straight.

```
train_index <- 49
n_total <- nrow(dataset[13:61, ])
dataset_train1 <- (dataset[13:(train_index), ]) # training 1-49
dataset_test <- dataset[(train_index + 1):61, ] # testing 49-61
predicted <- numeric((n_total + 12 - train_index)) #61-49=12
for (i in 1:(n_total - train_index + 12)) {
  dataset_train <- dataset[13:(train_index - 1 + i), ]
  arima.model <- auto.arima(ts(dataset_train$TotalInv))
  pred <- forecast(arima.model, 1)
  predicted[i] <- pred$mean
}
```

Next, I calculate error rates of the fitted values for Jan, 2017 to Jan, 2020 and output them into a table and on a plot:

```
## # A tibble: 6 x 4
##   Date      Actuals Predicted `Pred.Error Rate`
##   <date>      <dbl>    <dbl>          <dbl>
## 1 2019-08-01 4135940.  4112642.          0.01
## 2 2019-09-01 3532631.  3603197.          0.02
## 3 2019-10-01 4366414.  4200143.          0.04
## 4 2019-11-01 3326330.  3948279.          0.19
## 5 2019-12-01 3615882.  3471552.          0.04
## 6 2020-01-01 4331877.  4111353.          0.05
```



## 6. Adding Number of Locations as a Regressor to the ARIMA model

I considered to include one more variable into a model - *TotalUniqShipTos* - *Count of the Served Accounts*.

AIC of model with additional variable is 684.58, while AIC of the model without it is 672.30. Apparently, previous model is better.

I make a decision to move forward without adding this additional variable.

```
(arima2 <- auto.arima(ts1, D = 1, xreg = cbind(as.numeric(dataset$TotalUniqShipTos[13:49]))))
```

```
## Series: ts1
## Regression with ARIMA(0,0,0)(1,1,0)[12] errors
##
## Coefficients:
##      sar1      drift      xreg
##    -0.4589  20521.522  384.2810
## s.e.    0.2256   3911.615   84.2086
##
## sigma^2 estimated as 3.366e+10: log likelihood=-338.29
## AIC=684.58  AICc=686.58  BIC=689.45
```

```
arima1
```

```
## Series: ts1
## ARIMA(0,1,1)(0,1,0)[12]
##
## Coefficients:
##      ma1
```

```
##          -0.8114
## s.e.    0.1198
##
## sigma^2 estimated as 7.243e+10:  log likelihood=-334.15
## AIC=672.3   AICc=672.87   BIC=674.66
```

## 6. Forecasting 12 Months Sales (Defining Seasonality)


Previously, I predicted 12 months of validation set without considering seasons in the pattern. This time I am using AUTO.ARIMA algorithm with defined seasonality attribute (*seasonal = TRUE*). Also, I decide to use built-in **forecast** function instead of **for loop** approach.

```
ts_f <- ts(dataset$TotalInv[13:49], frequency = 12, start = c(2017,
1))
ts1_f <- window(ts_f)
summary(arima1_a <- auto.arima(ts1_f, D = 1, seasonal = TRUE))
```

```
## Series: ts1_f
## ARIMA(0,1,1)(0,1,0)[12]
##
## Coefficients:
##          ma1
##          -0.8114
## s.e.    0.1198
##
## sigma^2 estimated as 7.243e+10:  log likelihood=-334.15
## AIC=672.3   AICc=672.87   BIC=674.66
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
## Training set 26831.17 212183.7 126313.7 0.5024001 3.566076
##              MASE      ACF1
## Training set 0.3001183 -0.06986415
fcast2 <- forecast(arima1_a, h = 12)
```

The desired outcome is to have the black line (*observed values*) to stay within the shaded area (range of confidence levels). This indicates on a proper model forecasting technique while all predicted values are withing 80-95% of confidence level.

```
plot(fcast2)
lines((window(ts(dataset$TotalInv[13:61], frequency = 12, start = c(2017,
1)), D = 1)))
```



## 7. Prediction Error Rates Comparison Between Two Models

I calculate prediction error rates of the second model with defined seasonality as following:

```
actuals<-dataset[50:61,c(3,5)]

df_2_12mo <- cbind(Date = as.Date(actuals$Date, origin ="1970-01-01"),

                    actuals=actuals$TotalInv,
                    predicted=fcast2$mean[1:12],
                    error=round(((actuals$TotalInv-fcast2$mean[1:12])/actuals$TotalInv),2)

                    )
df_2_12mo<-data.frame(df_2_12mo)
df_2_12mo$Date = as.Date(df_2_12mo$Date, origin ="1970-01-01")
library(lubridate)
colnames(df_2_12mo)<-c("Date", "Actuals", "Predicted", "Pred.Error Rate")
df_2_12mo
```

##	Date	Actuals	Predicted	Pred.Error Rate
## 1	2020-02-01	4079161	3745981	0.08
## 2	2020-03-01	3554428	3945030	-0.11
## 3	2020-04-01	4537376	4443739	0.02
## 4	2020-05-01	3978675	4115515	-0.03
## 5	2020-06-01	3857241	4239019	-0.10
## 6	2020-07-01	4612485	4596395	0.00
## 7	2020-08-01	4214495	4667063	-0.11
## 8	2020-09-01	4286579	4063754	0.05
## 9	2020-10-01	5180160	4897537	0.05
## 10	2020-11-01	4459344	3857453	0.13
## 11	2020-12-01	4351108	4147005	0.05
## 12	2021-01-01	5059811	4863000	0.04

Now I recall outputs from the first model and second model to compare error rates. Magnitude of the error rates indicate that first model predicts better. However, I intend to consider two and discuss the output with the domain knowledge expert to have an insight what makes more sense. Variances described by error rates sometimes can be a result of rare business events that analysts are not aware of. So, having those variances would actually mean a correct prediction of the natural course of business.

```
df1 <- as.data.frame(tail(df_1_12mo, 12))
df1
```

##	Date	Predicted	Actuals	Pred.Error_Rate
## 1	2020-02-01	3535931	4079161	0.13
## 2	2020-03-01	4070634	3554428	0.15
## 3	2020-04-01	4174973	4537376	0.08
## 4	2020-05-01	3972652	3978675	0.00
## 5	2020-06-01	4037439	3857241	0.05
## 6	2020-07-01	4301916	4612485	0.07
## 7	2020-08-01	4011924	4214495	0.05
## 8	2020-09-01	4199740	4286579	0.02
## 9	2020-10-01	4574009	5180160	0.12
## 10	2020-11-01	4327525	4459344	0.03
## 11	2020-12-01	4680376	4351108	0.08
## 12	2021-01-01	5073783	5059811	0.00


df\_2\_12mo

##	Date	Actuals	Predicted	Pred.Error	Rate
## 1	2020-02-01	4079161	3745981	0.08	
## 2	2020-03-01	3554428	3945030	-0.11	
## 3	2020-04-01	4537376	4443739	0.02	
## 4	2020-05-01	3978675	4115515	-0.03	
## 5	2020-06-01	3857241	4239019	-0.10	
## 6	2020-07-01	4612485	4596395	0.00	
## 7	2020-08-01	4214495	4667063	-0.11	
## 8	2020-09-01	4286579	4063754	0.05	
## 9	2020-10-01	5180160	4897537	0.05	
## 10	2020-11-01	4459344	3857453	0.13	
## 11	2020-12-01	4351108	4147005	0.05	
## 12	2021-01-01	5059811	4863000	0.04	

```
a <- ggplot(df1, aes(x = Date, group = 1)) + geom_line(aes(y = Actuals/1e+06,
  col = "Actuals")) + geom_line(aes(y = Predicted/1e+06, col = "Predicted")) +
  xlab("Months") + ylab("Sales, in millions") + ggtitle("Actual vs Predicted Sales Values with Arima")

b <- ggplot(df_2_12mo, aes(x = Date, group = 1)) + geom_line(aes(y = Actuals/1e+06,
  col = "Actuals")) + geom_line(aes(y = Predicted/1e+06, col = "Predicted")) +
  xlab("Months") + ylab("Sales, in millions") + ggtitle("Actual vs Predicted Sales Values with Arima")

grid.arrange(a, b, nrow = 2)
```





## 8. 2021 Sales Forecast

Now I take all historical data from January 2017 to January 2021 and output the prediction for the next 12 months of 2021.

First, I proceed with utilizing the model by forecasting one month at a time, not counting for seasonality.

```
n_total <- nrow(dataset[13:61, ])  
dataset_train1 <- (dataset[13:(train_index), ]) # training 1-61  
predicted <- numeric(12)  
for (i in 1:(12)) {  
  dataset_train <- dataset[13:(60 + i), ]  
  arima.model <- auto.arima(ts(dataset_train$TotalInv))  
  pred <- forecast(arima.model, 1)  
  predicted[i] <- pred$mean  
}
```

Next, I utilize ARIMA model with defined seasonality and predict future values based on **forecast** built-in function:

```
ts_f2 <- ts(dataset$TotalInv[13:61], frequency = 12, start = c(2017,  
1))  
ts2_f <- window(ts_f2)  
summary(arima2 <- auto.arima(ts_f2, D = 1, seasonal = TRUE))
```

```
## Series: ts_f2  
## ARIMA(0,1,1)(0,1,0)[12]  
##  
## Coefficients:  
##          ma1  
##        -0.8122  
## s.e.      0.1219  
##  
## sigma^2 estimated as 8.416e+10: log likelihood=-503.92  
## AIC=1011.84   AICc=1012.2   BIC=1015.01  
##  
## Training set error measures:  
##              ME      RMSE      MAE      MPE      MAPE  
## Training set 35414.3 245180.5 164521 0.6166287 4.331813  
##              MASE      ACF1  
## Training set 0.348146 -0.0186137
```

```
fcast2021 <- forecast(arima2, h = 12)
```

Predicted sales for the rest of 2021 (February to December) from both models look like following. Side-by-side comparison shows that some months are nearly the same, however, there are a few differences that worth discussing with domain knowledge experts in Sales Department.

```
df_2021 <- tibble(Month_2021 = c(2:12), Predicted = c(predicted[1:11]),  
  Predicted_wSeas = fcast2021$mean[1:11])  
df_2021
```

```
## # A tibble: 11 x 3  
##   Month_2021 Predicted Predicted_wSeas  
##       <int>     <dbl>         <dbl>  
## 1         2  4536655.      4748728.  
## 2         3  4660806.      4223995.  
## 3         4  5013118.      5206944.
```

```
## 4          5 4682425.      4648243.
## 5          6 4867139.      4526808.
## 6          7 5022229.      5282053.
## 7          8 4842441.      4884063.
## 8          9 5016534.      4956146.
## 9         10 5075512.      5849727.
## 10        11 4998787.      5128911.
## 11        12 5138352.      5020675.
```

```
ggplot(df_2021, aes(x = as.factor(Month_2021), group = 1)) +
  geom_line(aes(y = Predicted/1e+06, col = "Predicted")) +
  geom_line(aes(y = Predicted_wSeas/1e+06, col = "Predicted w/Seasons")) +
  xlab("Months, 2021") + ylab("Sales, in millions") + ggtitle("2021 Predicted Sales. Two models comparison")
```



Considering the nature of business, I am in favor of the model with the defined seasonality. It is also worth looking at the values that lay within confidence intervals of 95% and 80%.

```
fcast2021
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Feb 2021	4748728	4376948	5120508	4180140	5317317
## Mar 2021	4223995	3845713	4602277	3645463	4802527
## Apr 2021	5206944	4822270	5591617	4618637	5795251
## May 2021	4648243	4257282	5039203	4050320	5246165
## Jun 2021	4526808	4129660	4923957	3919422	5134194
## Jul 2021	5282053	4878811	5685294	4665348	5898757
## Aug 2021	4884063	4474819	5293306	4258179	5509947
## Sep 2021	4956146	4540988	5371305	4321216	5591077
## Oct 2021	5849727	5428737	6270718	5205877	6493578
## Nov 2021	5128911	4702168	5555655	4476264	5781559
## Dec 2021	5020675	4588256	5453095	4359347	5682004
## Jan 2022	5729378	5291356	6167400	5059481	6399275